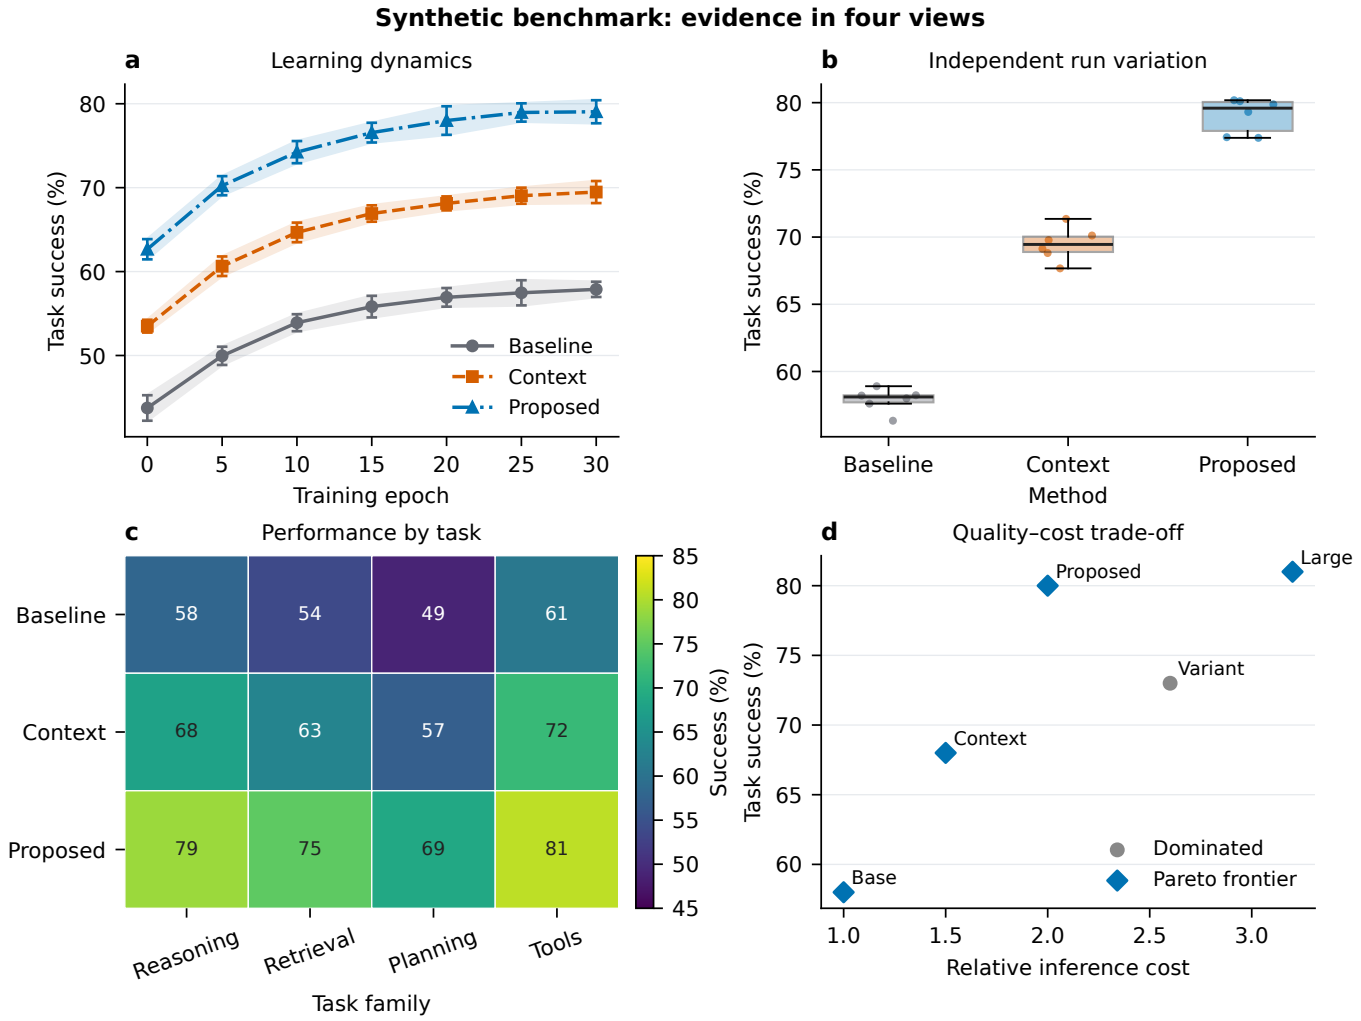
Synthetic benchmark: evidence in four views
a
Learning dynamics
Task success (%)
Training epoch
Baseline
Context
Proposed
b
Independent run variation
Task success (%)
Baseline
Context
Proposed
Method
c
Performance by task
Baseline
Context
Proposed
58
54
49
61
68
63
57
72
79
75
69
81
Reasoning
Retrieval
Planning
Tools
Task family
Success (%)
d
Quality–cost trade-off
Task success (%)
Proposed
Large
Variant
Context
Base
Dominated
Pareto frontier
Relative inference cost